

A brief introduction to Bayesian Statistics through Astronomical Applications (Lecture 3)


Cecilia Mateu J.
Instituto de Astronomía, UNAM, Ensenada

IA-UNAM, C.U.
30 de enero de 2015

More Examples

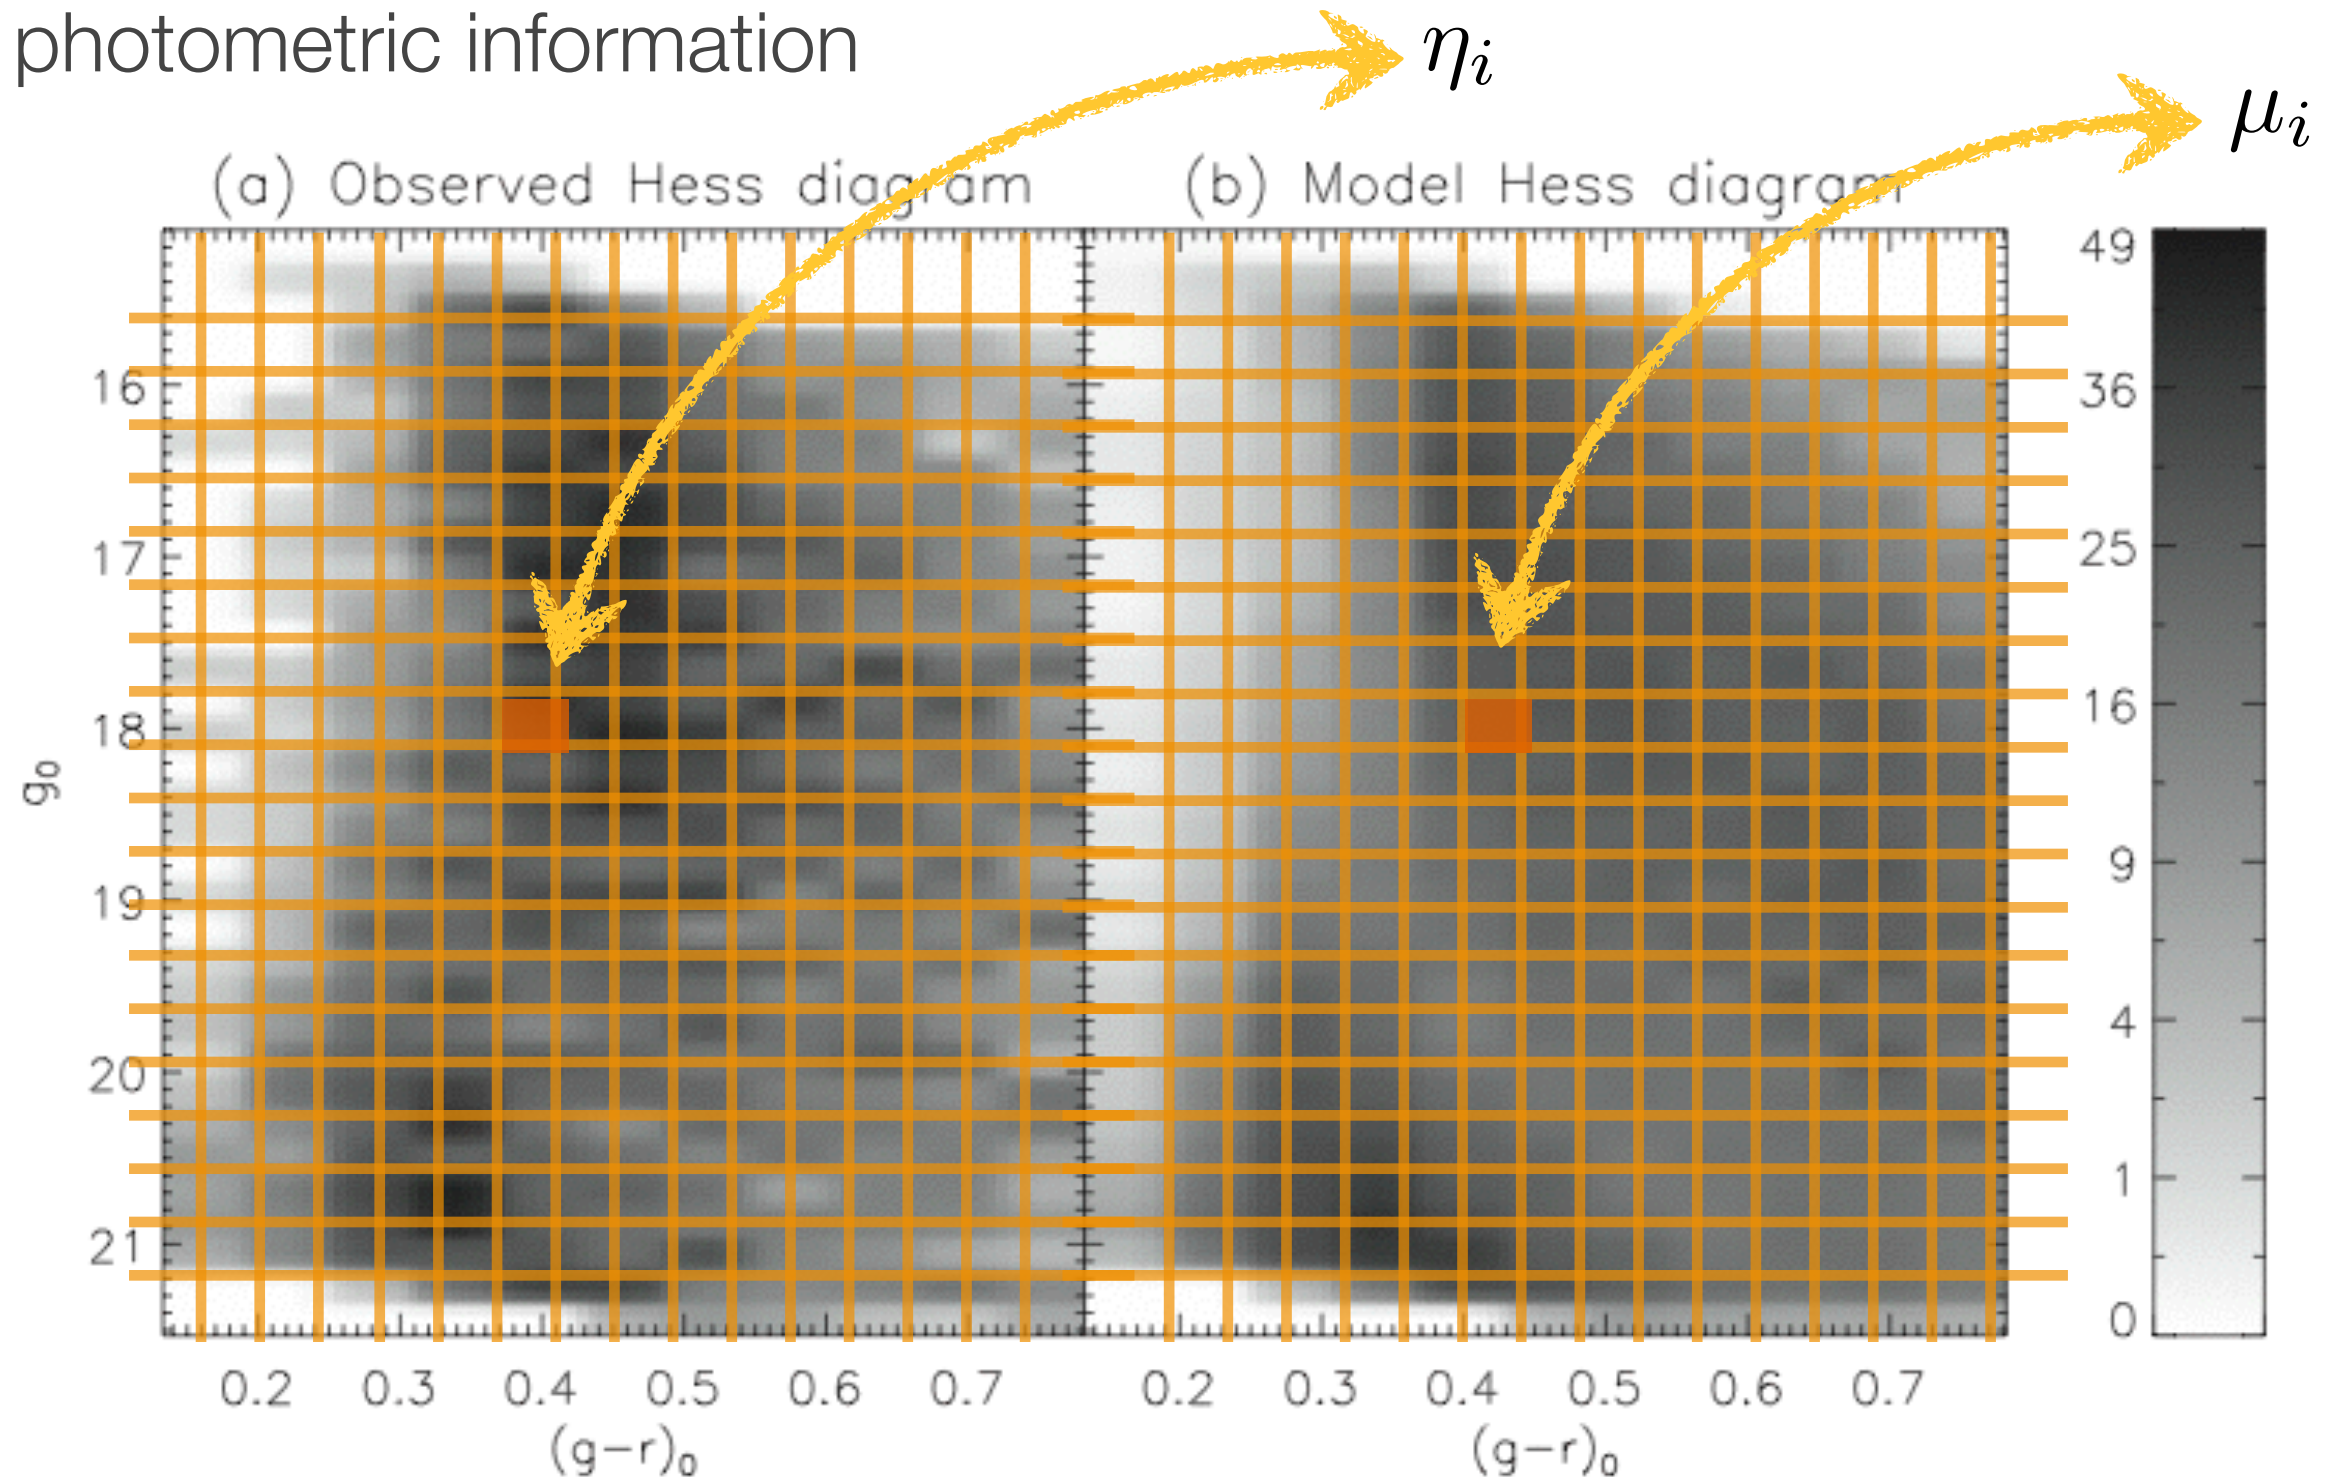
Fitting Density Profiles using Hess Diagrams

- The density profiles for the different components of the Galaxy determine the way a color-magnitude diagram looks for a given field in a particular line of sight
- We'd like to use this the other way around: from color-magnitude diagrams (CMDs) we want to *infer* the density profile parameters for the different Galactic components



Fitting Density Profiles using Hess Diagrams

- We can use a procedure similar to what we saw yesterday, but with photometric information
- Lets make a grid in the observed and model CMDs



Fitting Density Profiles using Hess Diagrams

- ◆ In general our free parameters are:

$$\vec{\theta} = (h_z^{TnD}, h_R^{TnD}, C_{TnD}, h_z^{TkD}, h_R^{TkD}, C_{TkD}, n, q, C_H)$$

- ◆ The likelihood function is again Poissonian

$$L \equiv p(\{\eta\}|\vec{\theta}) = \prod_{i \in V_S} p(\eta_i|\vec{\theta}) = \prod_{i \in V_S} \frac{\mu_i^{\eta_i} e^{-\mu_i}}{\eta_i!}$$

$$\ln L = \sum_i \eta_i \ln \mu_i - \mu_i$$

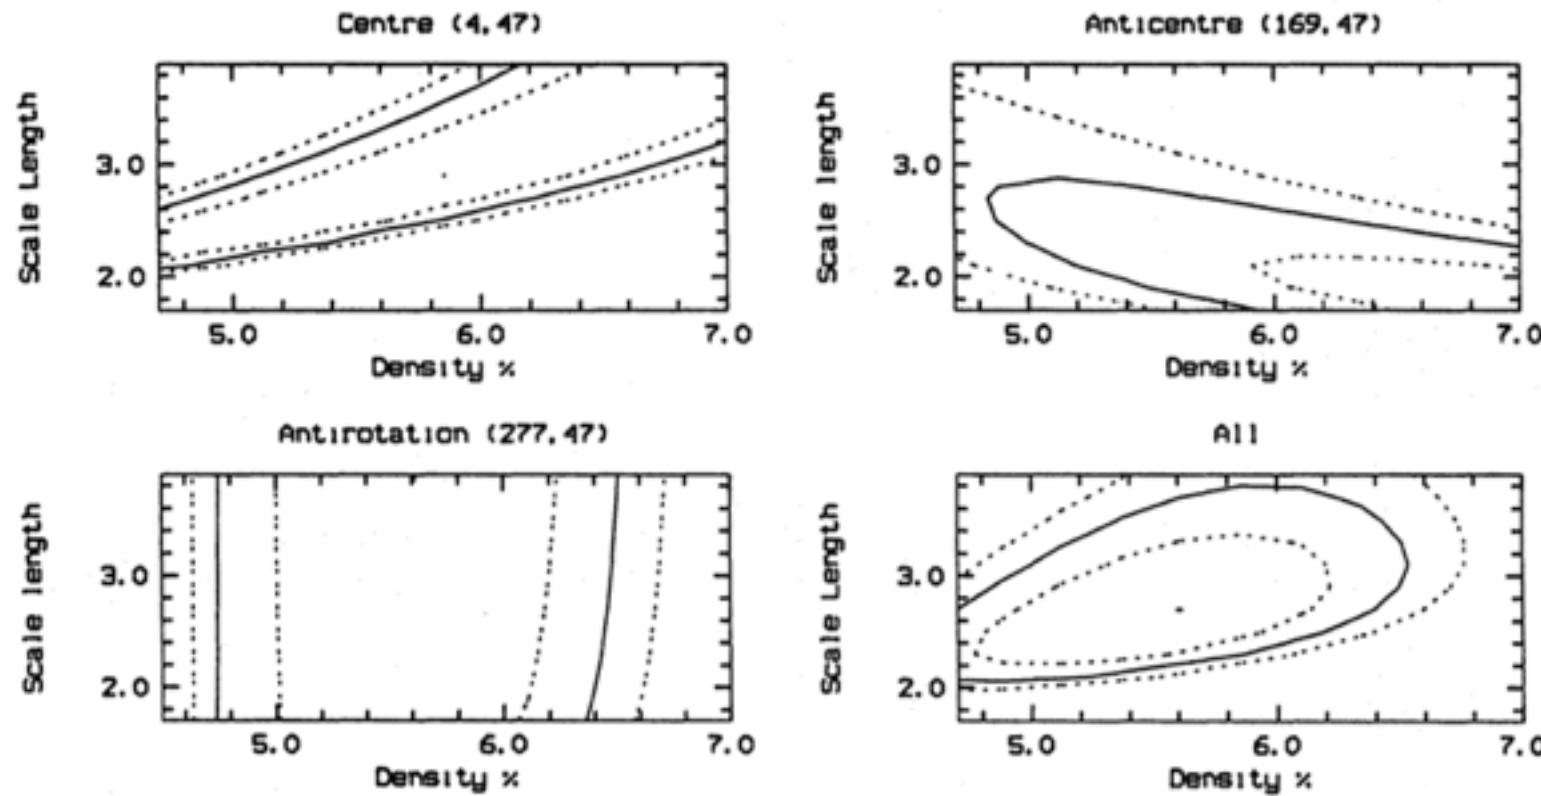

- ◆ This is commonly used in the literature to fit density profiles using different types of photometric data: e.g. de Jong et al. (2010), Robin et al. (1996)

η_i - *observed* number or RRLS
on i -th bin

μ_i - *predicted* number or RRLS
on i -th bin

Density Profiles from Hess Diagram fitting

- Robin et al. (1996) show the result of fitting separately using CMDs from fields in different lines of sight and then combining them all




- de Jong et al. (2010) use SEGUE data which spans a large range of longitudes and crosses the Galactic disk
- They explore the effect of removing particular fields with known substructure, such as the Sgr and Monoceros streams

Exploring the parameter space in high dimensionality problems: Markov Chain Monte Carlo


Computing the Posterior in many-parameter problems

- As we have already discussed, the Posterior can only be found by direct evaluation in problems with very few parameters ($< 3 - 4$)
- We would like to have a way of exploring the parameter space efficiently, spending more computation time around high-probability areas than around low probability ones




Obtaining samples from the Posterior

- One way is to try to obtain Posterior samples, i.e. a random realisation of the Posterior PDF
- If one has a random realisation of the Posterior with N samples, the Posterior is simply the N -dimensional histogram of this samples
- Having posterior samples, Marginalization is trivial, just the histogram in any lower number of dimensions is the marginal posterior!
- Uncertainties can be easily computed as the standard deviation or percentiles in the resulting histograms




Posterior Samples: Von Neumann Rejection

- A simple way of doing this is with the Von Neumann Rejection Technique, also known as Accept-Reject:
 - Generate random uniformly-distributed samples (x,y) with $x_o < x < x_f$ and $0 < y < \max(\text{Posterior})$
 - Accept only the samples for which $y < \text{Posterior}(x)$
 - ... that's it, the accepted points are distributed as the posterior
- This is quite simple and works in any number of dimensions!
- However...



Posterior samples: Markov Chain Monte Carlo

- Von Neumann rejection can still be very inefficient⁺ for most problems, so Markov Chain Monte Carlo (MCMC) is preferred
- The idea of MCMC is to start from a point and explore the parameter space by taking steps that may be accepted or rejected, such that the Markov chain:
 - Tends to walk towards higher probability areas
 - Tends to avoid low probability areas




(from Tom Loredó's Lecture Notes)

- Note that the samples are not completely independent, there is some correlation
- After a while, the chain 'forgets' the initial conditions and the accepted (independent) samples have a PDF that is proportional to the Posterior

A two-parameter problem

- We observe the following distribution of N pairs (x_i, y_i)
- It seems reasonable to assume they were drawn from a random distribution, so let's use a gaussian model with known $\sigma_x = \sigma_y = 1$ and μ_x, μ_y the unknown means in the X and Y directions
- The likelihood is expressed as

$$P(\{x_i, y_i\} | \mu_X, \mu_Y) = \prod_{i=1}^N e^{\frac{1}{2}[(x_i - \mu_X)^2 + (y_i - \mu_Y)^2]}$$



- Assuming a uniform prior probability for μ_x, μ_y , the posterior is therefore given by

$$P(\mu_X, \mu_Y | \{x_i, y_i\}) = \prod_{i=1}^N e^{\frac{1}{2}[(x_i - \mu_X)^2 + (y_i - \mu_Y)^2]}$$



MCMC: The Metropolis-Hastings Recipe

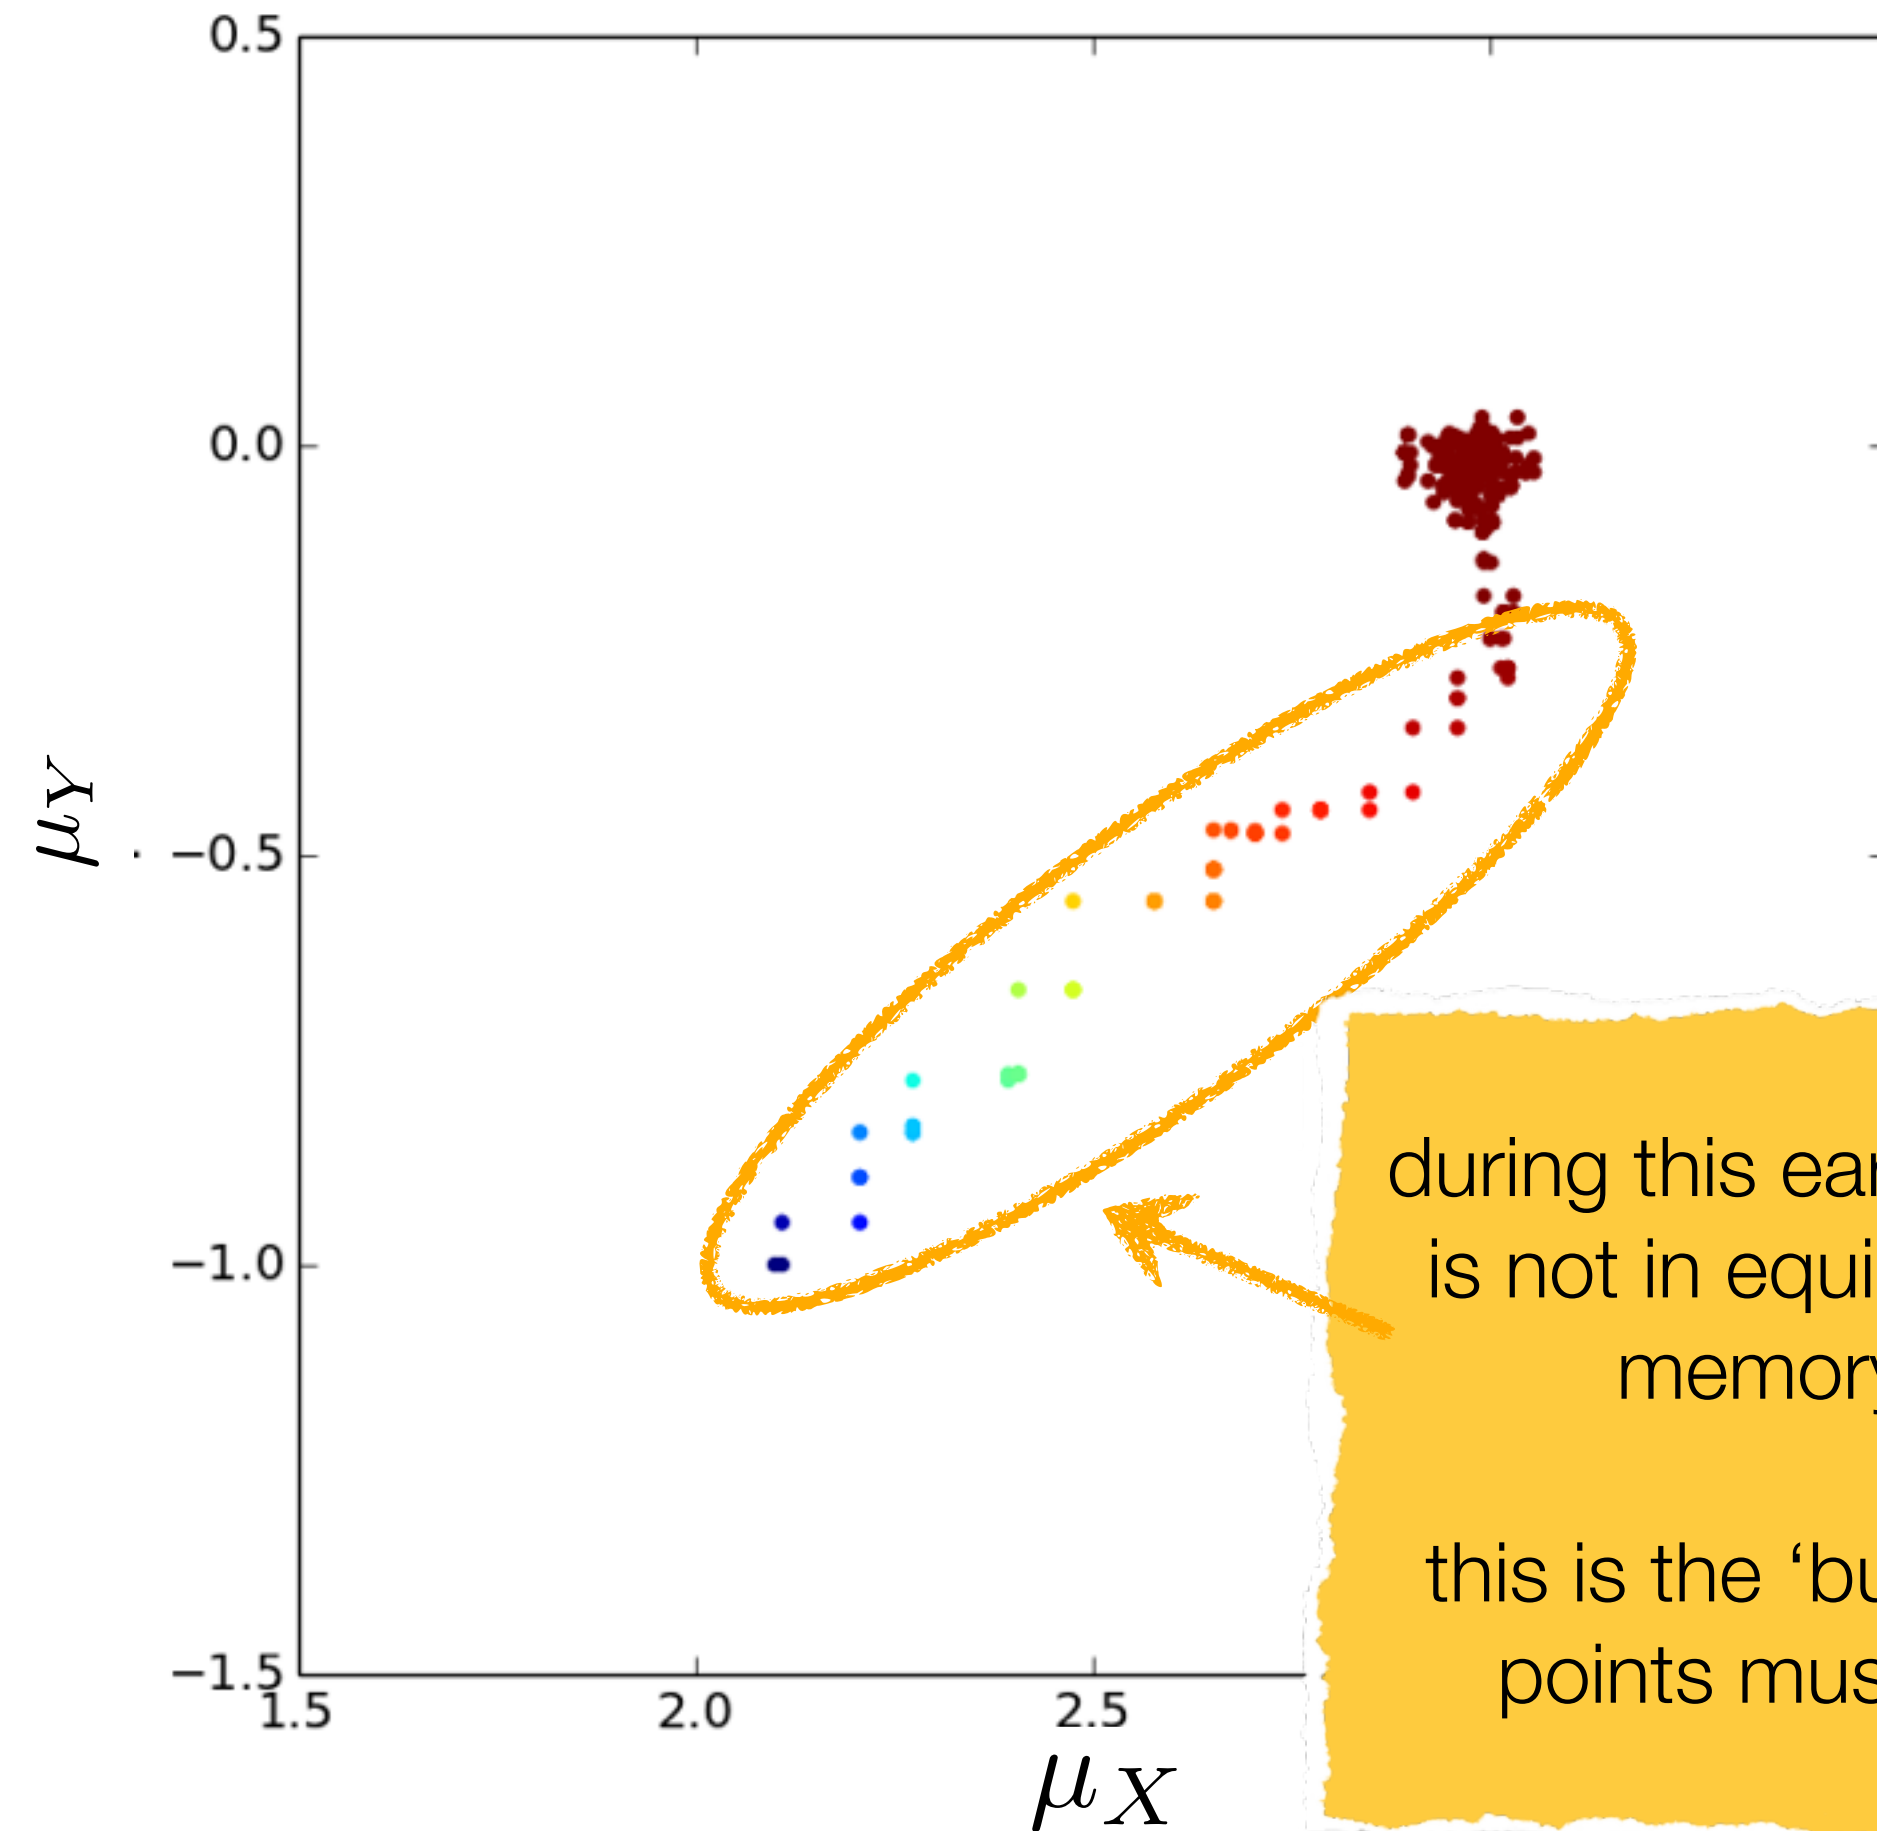
from Hogg et al. 2010

- 1-Choose an initial position for the model params $\{\mu_i\}$
- 2-Advance a step (in one parameter) randomly $\rightarrow \mu_{i+1}$
- 3-Evaluate the posterior at current position $P(\mu_{i+1})$
- 4-Draw a random number R with uniform probability in the range $0 < R < 1$
 - If $R < P(\mu_{i+1})/P(\mu_i)$, keep the point and add it to the chain
 - if not, go back to the previous step and re-add it to the chain
 - repeat ...
- **the set of $\{\mu_i\}$ obtained is a random realization of the Posterior !**

MCMC: The Metropolis-Hastings Recipe

- The step size must be chosen so that the acceptance fraction (fraction of points accepted in the chain) lies between ~ 0.2 and ~ 0.5 (see Hogg et. al. 2010 and Foreman-Mackey et al. 2013)
- This algorithm is a piece of cake to write, excellent for playing around to develop some intuition as to how the MCMC works
- The problem is that fine-tuning the chain when the number of parameters is large is highly non-trivial! (there's no way of guessing it a priori)
- This is solved by MCMC implementations like **emcee** (in Python, Foreman-Mackey et al. 2013) that use algorithms more sophisticated than Metropolis-Hastings, with very few free parameters







during this early stage the chain
is not in equilibrium yet, it has
memory of its path

this is the 'burn-in' stage, this
points must be discarded




A two-parameter problem

- MCMC sampling



MCMC samples

- The set of points obtained in the final Markov Chain is a random realization of the Posterior PDF
- The mode of the histogram gives the most probable value of each parameter
- The percentiles give the credible regions



More Suggested Bibliography

- Hogg, Bovy & Lang (2010)

Data analysis recipes: Fitting a model to data*

David W. Hogg

*Center for Cosmology and Particle Physics, Department of Physics, New York University
Max-Planck-Institut für Astronomie, Heidelberg*

Jo Bovy

Center for Cosmology and Particle Physics, Department of Physics, New York University

Dustin Lang

*Department of Computer Science, University of Toronto
Princeton University Observatory*

Approximate Bayesian Computation (ABC)

Approximate Bayesian Computation (ABC)

- ◆ Option for cases where there's no analytic likelihood, but there is enough knowledge about the problem to do forward modelling

Basic ABC algorithm

For the observed data $y_{1:n}$, prior $\pi(\theta)$ and distance function ρ :

Algorithm*

- 1 Sample θ^* from prior $\pi(\theta)$
- 2 Generate $x_{1:n}$ from forward process $f(y \mid \theta^*)$
- 3 Accept θ^* if $\rho(y_{1:n}, x_{1:n}) < \epsilon$
- 4 Return to step 1

Generates a sample from an approximation of the posterior:

$$f(x_{1:n} \mid \rho(y_{1:n}, x_{1:n}, \theta) < \epsilon) \cdot \pi(\theta) \approx f(y_{1:n} \mid \theta) \pi(\theta) \propto \pi(\theta \mid y_{1:n})$$

*Introduced in Pritchard et al. (1999) (population genetics)

"Though there be no such thing as
Chance in the world; our ignorance
of the real cause of any event has
the same influence on the
understanding"

-David Hume (1748)